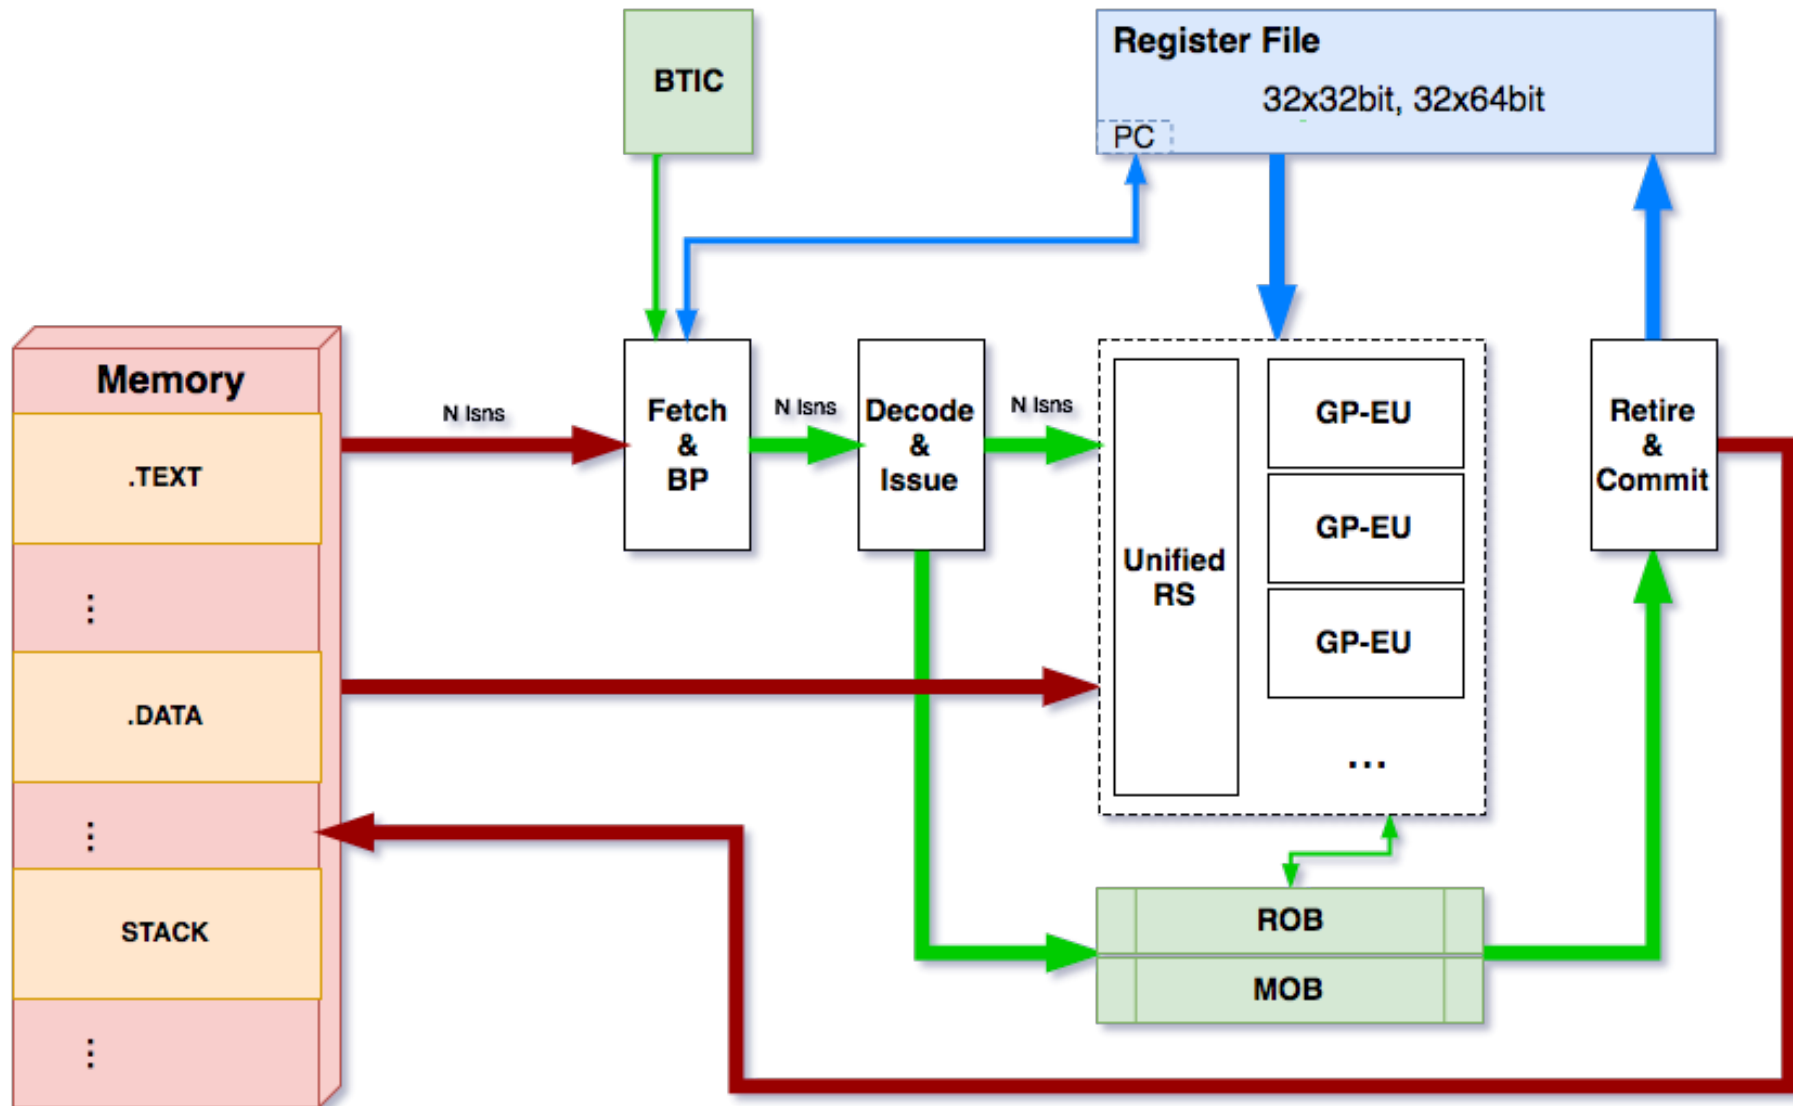

Kraken

Joy Yeh
cy13308



Avg. IPC: 3.9
Max IPC: 6

Features 1

- **Pipelined:**
 - 4 major stages: Fetch, Decode, Execute, Commit
 - Sub-pipelines
- **Superscalar:**
 - OoO Execute, In-Order completion, arbitrary number of execution units (assumed GP).
 - Speculative Execution – arbitrary level*
- **Branch Prediction:**
 - Fixed, Static, Dynamic n-bits
 - Branch Target Instruction Cache (BTIC) with Relative Instruction Pointers (RIPs)

Features 2

- **Dependency Reduction:**
 - Register file “extension” using ROB
 - Second ROB (Memory Order Buffer) for load/stores*
- **ISA:**
 - AArch64, Neon, VFP (ARM vixl decoding library)
 - Executes real ARM64 benchmark binaries compiled from C. Can't link libraries yet, just pure C with a core Maths library

Latency Model

- | | |
|--------------------------|-----|
| • ADD/SUB/SHIFT/MOV... | 1 |
| • Shifted ADD/SUB | 2 |
| • Load/Store (assume L1) | 4/5 |
| • Mul | 4 |
| • Div | 20 |
| • Branches | 5 |

(Data pieced together from different processors...)

Benchmarks

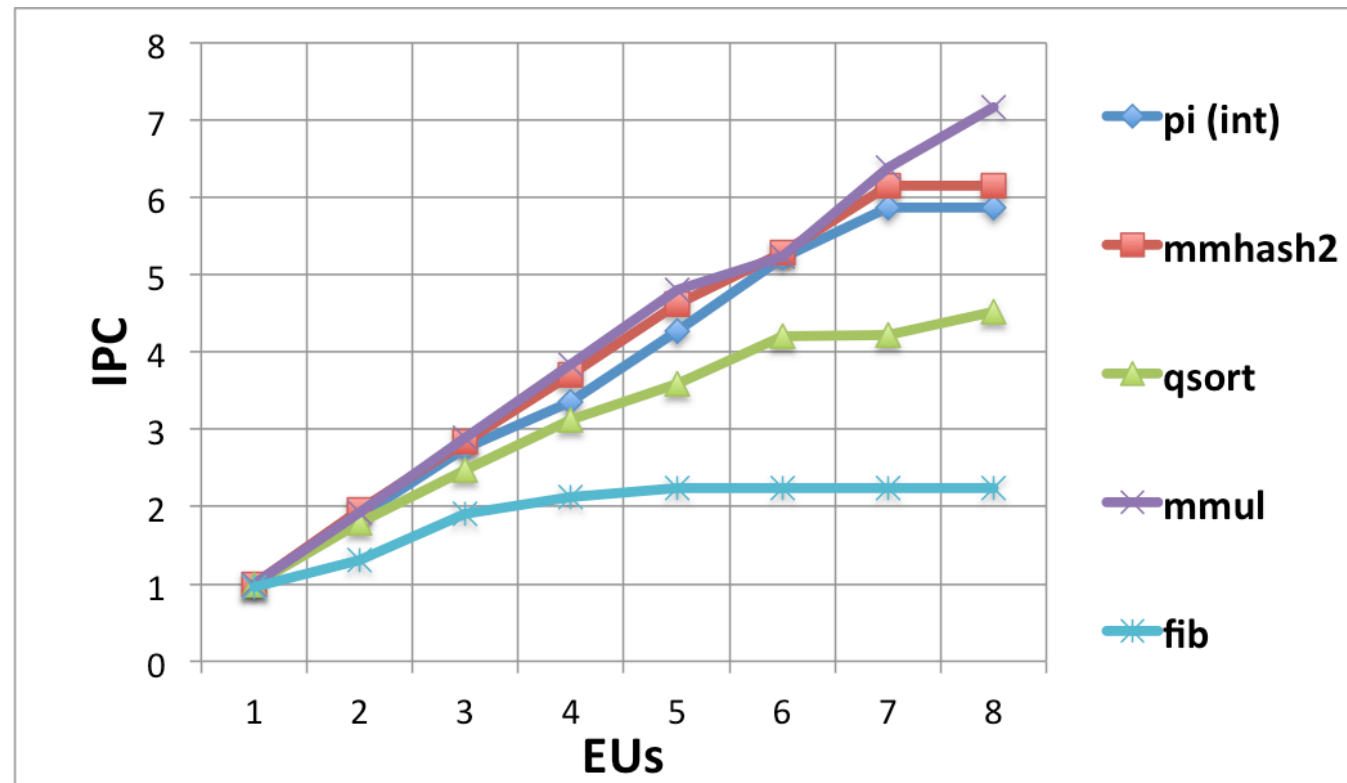
Ranging ($10^{(6\sim 8)}$ instructions):

- Unrolled: Fac
- Tight loops: Fac, Fibs
- Loaded loops: QSort, Pi, MMul, BUDE
- Recursion: Fac
- Artificial/Misc.: MMHash2, Nops

H1

IPC increases as the number of general purpose execution units, with diminishing returns

Config:
6 EU
5-bit BP
OOOSS
Clang -O0



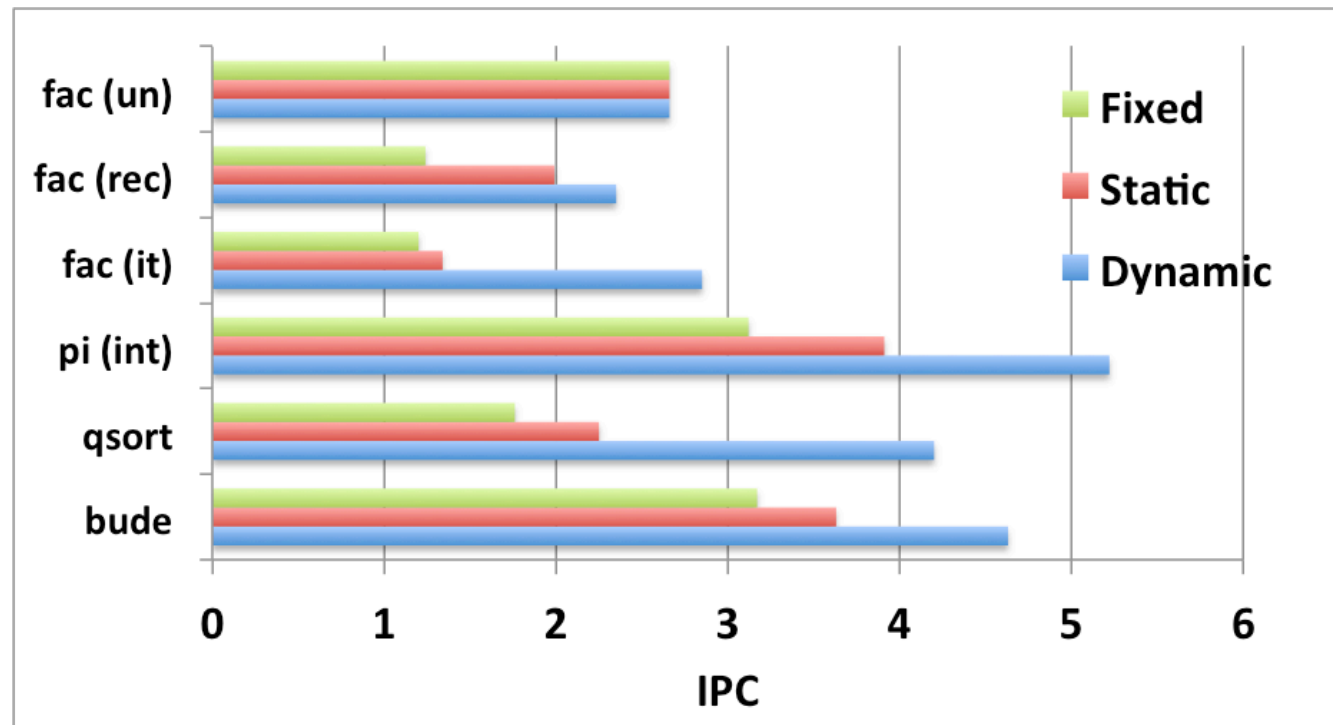
Observations

- Matrix multiply is very easily paralised hence the continued trend, but other codes have more intra-loop dependencies so they stop around 6
- Really tight loops like Fib suffer a lower plateau. Pi, hash and MMul have a lot more work per loop to parallelise. **Result differs depending on code generation method.

H2

Ideally, branch prediction should deliver an IPC almost equal to unrolled loop

Config:
6 EU
5-bit BP
OOOSS
Clang -O0



Observations

- Dynamic Factorial gives HIGHER IPC than unrolled Factorial!
 - Misc. instructions in the loop + speculation allowed extra instruction counts
- Large amount of work within each loop reduces the advantage of BP (BUDE - less diff.).


H3

Compiler optimisation should help increase IPC... or not?

Config:

6 EU
5-bit BP
OOOSS

Clang -O3
-ffast-math
-fvectorize
-mcpu=cortex-
a57



Observations

- Compiler optimised for code density and ultimately minimum cycle count, and not IPC (MMHash2 and Qsort have lower IPC in O3).
- The total cycle count is a better/more important measure when considering real world application. Lower IPC can be improved by adding SMT.